

Deep Learning with Point Clouds

Romain Thalineau

qwertee.io

romain.thalineau@qwertee.io

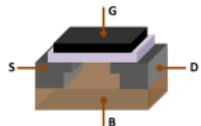


@romaintha

PyData Cluj-Napoca, February 20, 2019

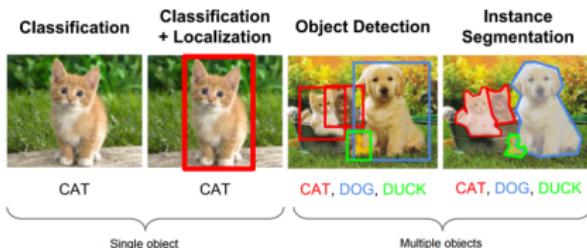
About

- Ph.D. in Physics (Grenoble, France), working on single electron spin qubits
- Worked at Infineon (Munich, Germany) as a data scientist.
- Living in Cluj since 2016 and now working as a machine learning engineer & data scientist at Qwertee.



From 2D to 3D perception

- Impressive developments over the last decade in the field of 2D vision.



Credit: Kdnuggets

- Advanced robotic systems require visual perception capabilities beyond 2D images.



Overview

1 Point Clouds

- Definition
- Visualization
- Learning tasks

2 Learning features from point clouds

- Learning features from regular data
- Point cloud problem statement

3 PointNet

- PointNet proposal
- PointNet architecture
- PointNet PyTorch implementation
- Experiments

4 Beyond PointNet

- PointNet++
- PointCNN

Overview

1 Point Clouds

- Definition
- Visualization
- Learning tasks

2 Learning features from point clouds

- Learning features from regular data
- Point cloud problem statement

3 PointNet

- PointNet proposal
- PointNet architecture
- PointNet PyTorch implementation
- Experiments

4 Beyond PointNet

- PointNet++
- PointCNN

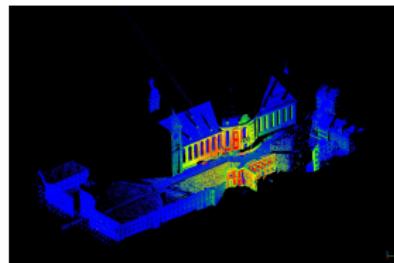
Point Clouds

- A point cloud is just a set of points. E.g.:

x	y	z
0.07388	0.16975	-0.19326
0.15286	0.15050	0.24355
0.20948	0.15050	0.29081
...

- By nature, point clouds are **irregular** (density) and **unordered**
⇒ invariant to permutations
- A point cloud may also hold additional features such as RGB or normals.

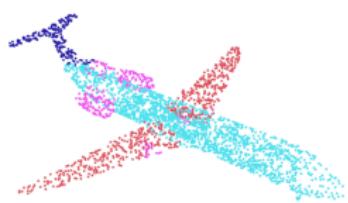
Example of point clouds



Semantic3D dataset



S3DIS dataset



ShapeNet dataset

Visualizing point clouds

- Visualize point clouds with Python

```
1 import open3d  
2  
3 point_cloud = open3d.read_point_cloud('pc_file.txt')  
4 open3d.draw_geometries([point_cloud])
```

- Use an OSS such as CloudCompare

Learning tasks

Classification¹



mug?

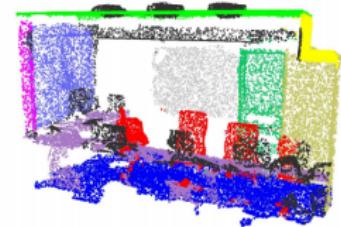


table?



car?

Scene/object semantic segmentation¹



¹Source: PointNet

Overview

1 Point Clouds

- Definition
- Visualization
- Learning tasks

2 Learning features from point clouds

- Learning features from regular data
- Point cloud problem statement

3 PointNet

- PointNet proposal
- PointNet architecture
- PointNet PyTorch implementation
- Experiments

4 Beyond PointNet

- PointNet++
- PointCNN

Learning features from regular data

input

3	0	0	0	0	0	1
0	1	1	3	1	0	0
2	0	2	0	0	0	0
1	3	1	1	3	0	0
0	0	0	0	2	0	0
2	1	3	2	0	0	0
1	3	2	1	0	0	0

conv

-1	-2	-1
0	0	0
1	2	1

1	4	2	-1
5	0	-2	2
-4	-4	0	4
-1	3	1	-5

activation
(ReLU)

pooling

1	4	2	0
5	0	0	2
0	0	0	4
0	3	1	0

Detailed description: The diagram illustrates a max pooling step. It starts with a 4x4 input grid and a 2x2 kernel. The kernel slides over the input with a stride of 2. The max values in each kernel's receptive field are highlighted with a green box. These max values are then mapped to a 2x2 output grid. The top-left cell of the output grid contains the value 5, and the bottom-right cell contains the value 3, indicating that the max value from the bottom-left input cell (0) was not selected because it was not the maximum in its local 2x2 window.

- Use convolution to leverage 3 important idea: **sparse connectivity, parameter sharing, equivariance to translation**
- Use max pooling for \sim **invariance to translation**.

Learning features from regular data

- Stacking simple blocks (CONV→RELU→POOL) allows for learning higher level features¹

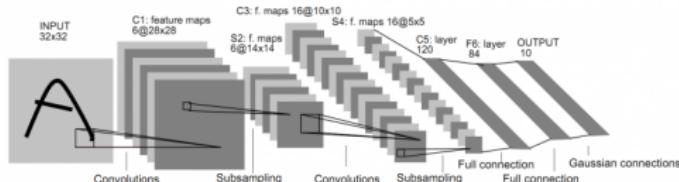
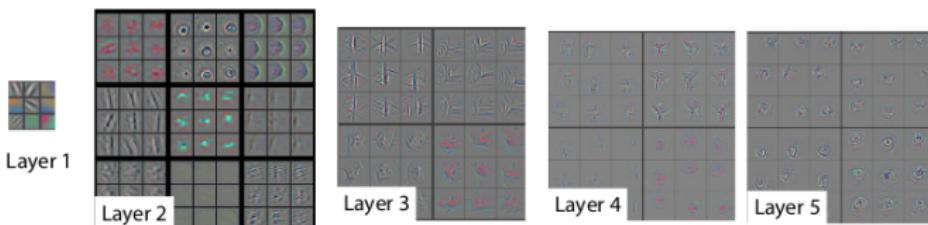


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

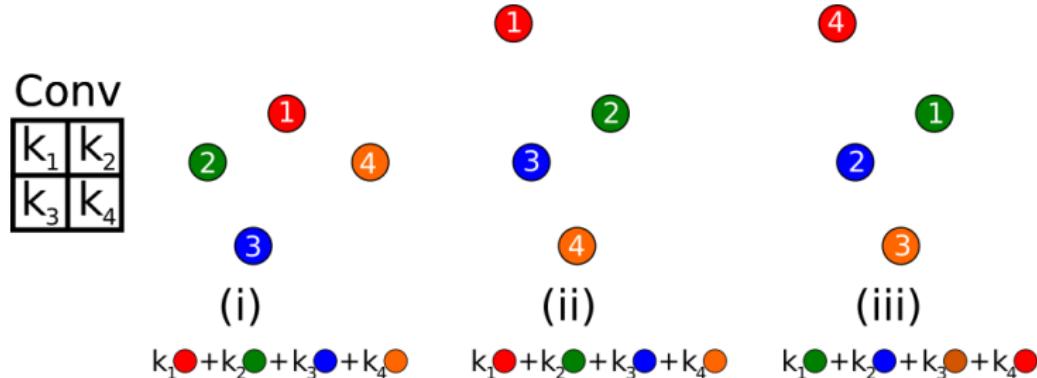
- Visualizing convolutional networks²



¹ Gradient based learning applied to document recognition, Lecun et al

² Visualizing and Understanding Convolutional Networks, Zeiler & Fergus

Point cloud problem statement



- Applying convolution kernel to point cloud:
 - description of shape $\text{conv}(i) = \text{conv}(ii)$
 - variance to ordering $\text{conv}(ii) \neq \text{conv}(iii)$
- We need functions f such that :
 - $f(i) \neq f(ii) \rightarrow$ Capturing local structures
 - $f(ii) = f(iii) \rightarrow$ Invariant to permutations
 - Invariant to geometric transformations such as rotations or translations of the whole point set .

Overview

1 Point Clouds

- Definition
- Visualization
- Learning tasks

2 Learning features from point clouds

- Learning features from regular data
- Point cloud problem statement

3 PointNet

- PointNet proposal
- PointNet architecture
- PointNet PyTorch implementation
- Experiments

4 Beyond PointNet

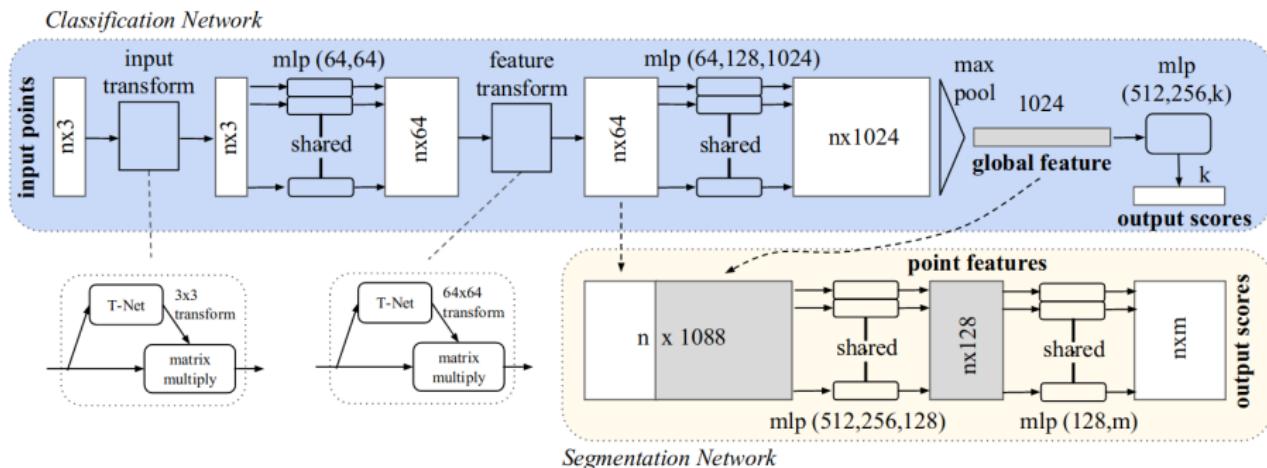
- PointNet++
- PointCNN

PointNet proposal

- Permutation invariance: by using a max pooling function.
 $f(x_1, \dots, x_n) = g(h(x_1), \dots, h(x_n))$ where:
 - h is a MLP network.
 - g is the max pooling function.
- Invariance to geometric transformations (rotations, translations...): by learning a transformation network¹
- By its design, PointNet does not capture well the local structure induced by the distance metric.

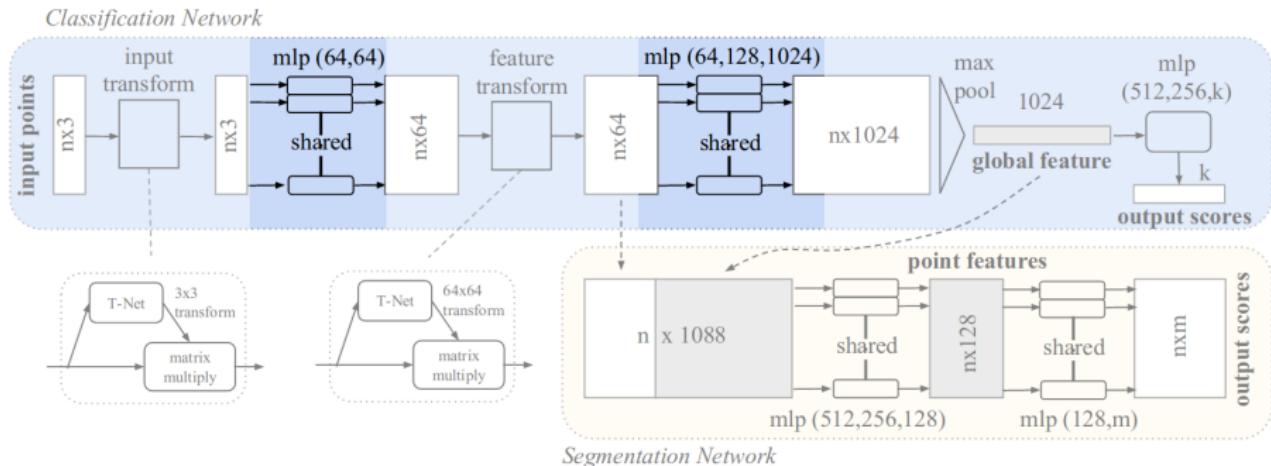
¹Slightly similar to "Spatial transformer network" by Jaderberg et al

PointNet architecture



Reference: "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", Qi et al

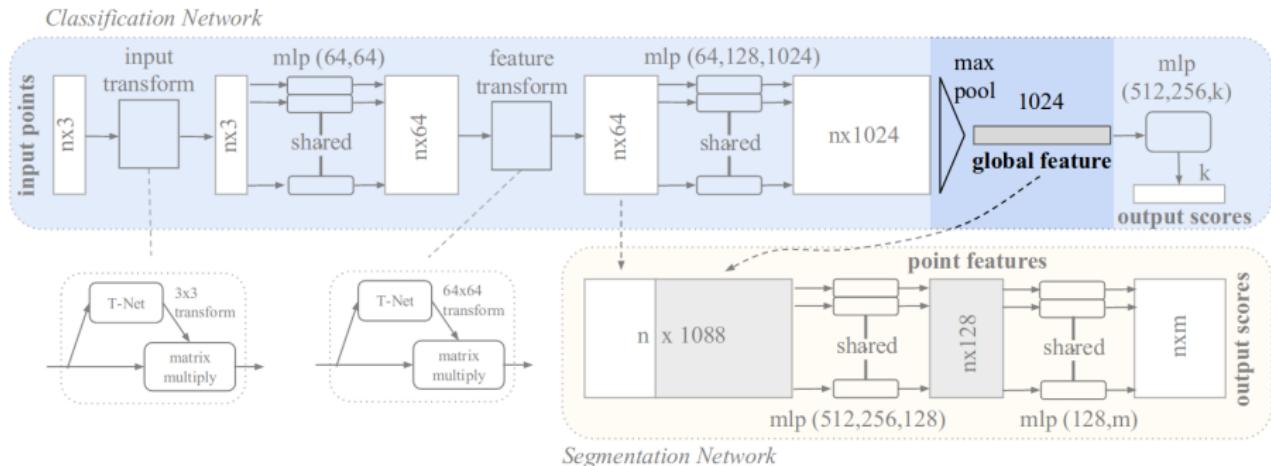
PointNet architecture



Shared MLP to learn spatial encoding of each point.

Note: Shared MLP is equivalent to a 1d conv with kernel size of 1

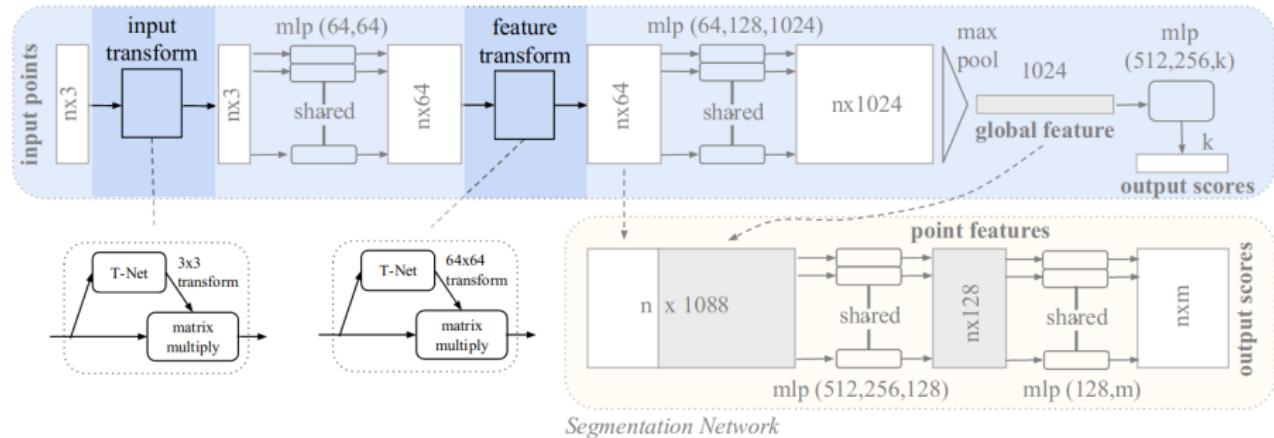
PointNet architecture



Symmetry function (max pool) in order to make the model invariant to input permutations.

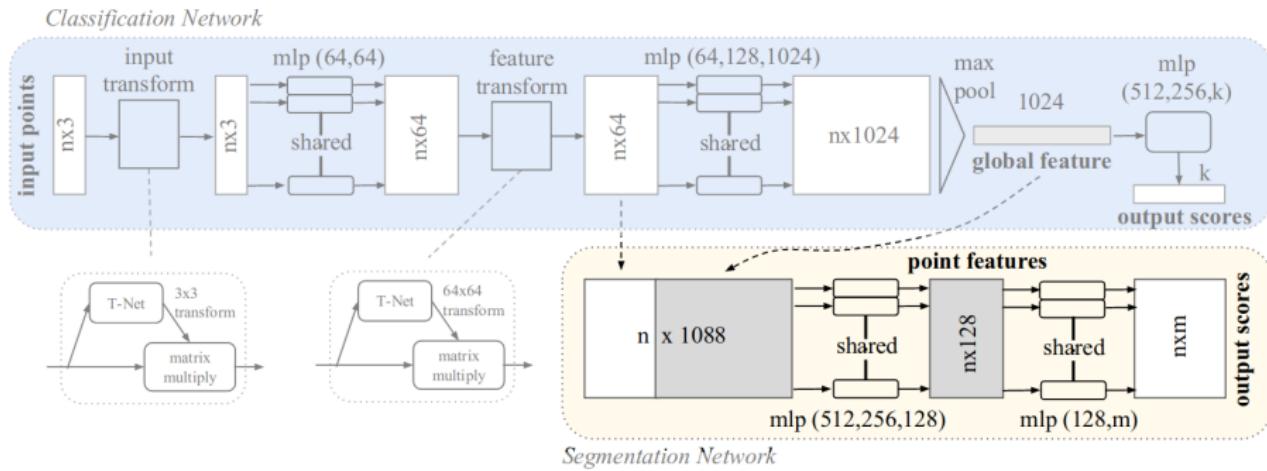
PointNet architecture

Classification Network



Joint alignment network to learn invariance to some geometric transformations. Similar to "Spatial Network Transformer", Jaderberg et al

PointNet architecture



PyTorch 101

- Pytorch tensors similar (\sim API) to numpy ND-arrays, but can be used on GPU.

```
1 import torch  
2  
3 x = torch.tensor([1, 2])  
4 x = torch.rand(2, 2)  
5 x.shape  
6 y = x + 2
```

- Automatic differentiation

```
1 x = torch.rand(2, 2)  
2 y = x + 2  
3 z = y.mean()  
4  
5 z.backward()  
6 x.grad
```

PyTorch 101

- Define a network

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class SimpleCNN(nn.Module):
6
7     def __init__(self):
8         super(SimpleCNN, self).__init__()
9         self.conv = nn.Conv2D(3, 64, 3)
10        self.fc = nn.Linear(64*32*32, 10)
11
12    def forward(self, x):
13        x = self.conv(x)
14        x = F.relu(x)
15        x = F.max_pool2d(x)
16        x = x.view(x.shape[0], -1) # flatten
17        x = self.fc(x)
18
19        return x
```

PyTorch 101

- Train a network

```
1 import torch.nn as nn
2 import torch.optim as optim
3
4 model = SimpleCNN()
5 criterion = nn.CrossEntropyLoss()
6
7 optimizer = optim.SGD(model.parameters(), lr=0.01)
8
9 for data, target in train_data_generator:
10     optimizer.zero_grad()
11     model = model.train()
12     output = model(data)
13     loss = criterion(output, target)
14     loss.backward()
15     optimizer.step()      # Does the update
```

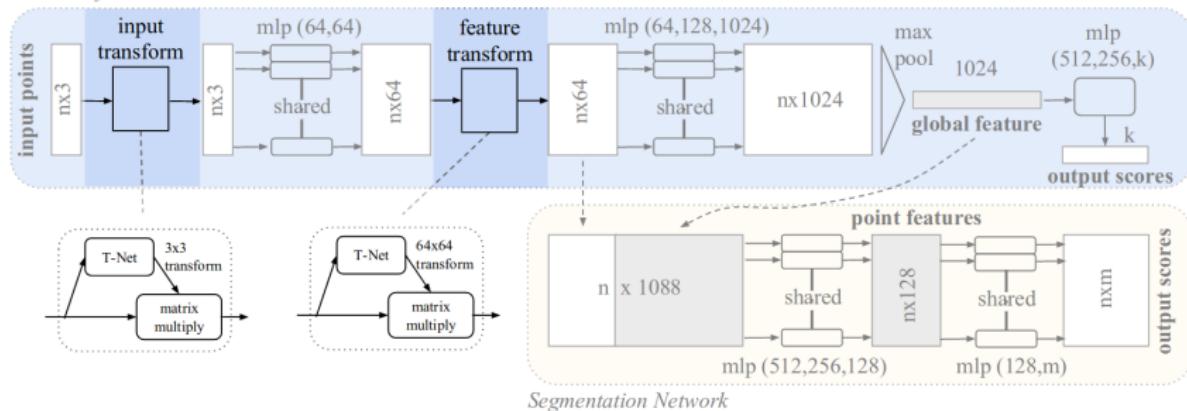
PointNet PyTorch: Transformer

Full implementation on my GitHub account

Warning: The implementation code is simplified for the presentation sake.

• Transformer Network

Classification Network

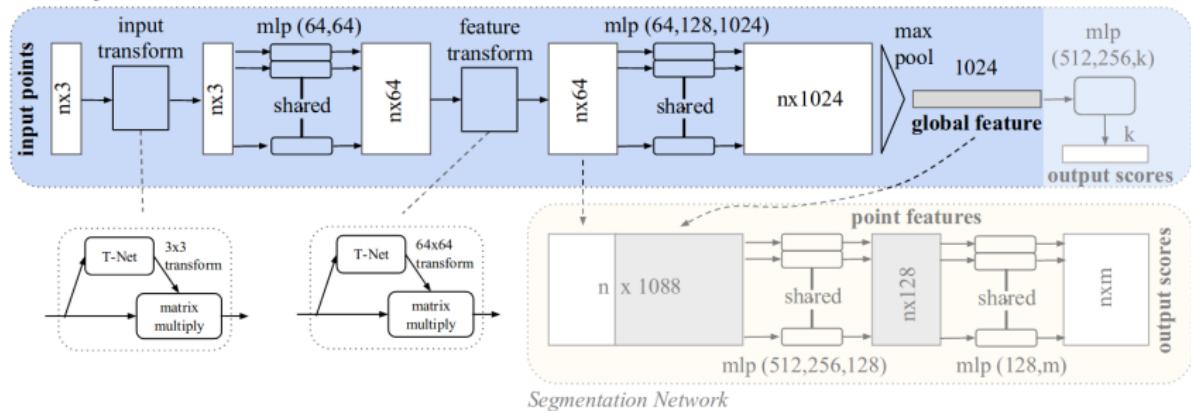


PointNet PyTorch: Transformer

```
1 class TransformationNet(nn.Module):  
2  
3     def forward(self, x): # x.shape = (B, N, C)  
4         num_points = x.shape[1]  
5         x = x.transpose(2, 1) # (B, C, N)  
6         x = F.relu(nn.Conv1d(C, 64, 1)(x)) # (B, 64, N)  
7         x = F.relu(nn.Conv1d(64, 128, 1)(x)) # (B, 128, N)  
8         x = F.relu(nn.Conv1d(128, 1024, 1)(x)) # (B, 1024, N)  
9  
10        x = nn.MaxPool1d(num_points)(x) # (B, 1024, 1)  
11        x = x.view(-1, 1024) # (B, 1024)  
12  
13        x = F.relu(nn.Linear(1024, 512)(x)) # (B, 512)  
14        x = F.relu(nn.Linear(512, 256)(x)) # (B, 256)  
15        x = nn.Linear(256, output_dim^2)(x) # (B, output_dim^2)  
16  
17        identity_matrix = torch.eye(output_dim)  
18        x = x.view(-1, output_dim, output_dim)  
19        return x + identity_matrix # (B, output_dim, output_dim)
```

PointNet PyTorch: Base network

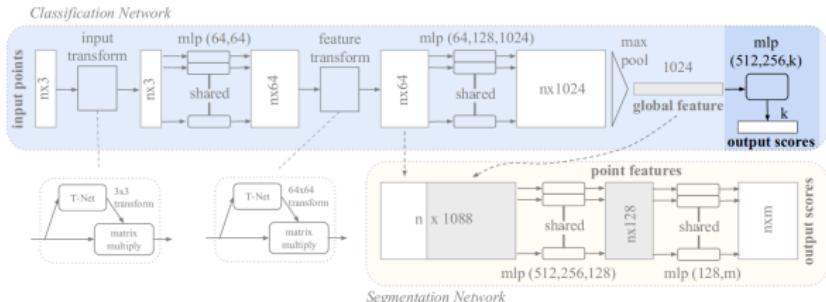
- Base Network
Classification Network



PointNet PyTorch: Base

```
1  class BasePointNet(nn.Module):  
2  
3      def forward(self, x): # x.shape = (B, N, C)  
4          num_points = x.shape[1]  
5          input_transform = TransformationNet(C)(x)      # (B, C, C)  
6          x = torch.bmm(x, input_transform)             # (B, N, C)  
7  
8          x = x.transpose(2, 1)                         # (B, C, N)  
9          x = F.relu(nn.Conv1d(C, 64, 1)(x))           # (B, 64, N)  
10         x = F.relu(nn.Conv1d(64, 64, 1)(x)))        # (B, 64, N)  
11         x = x.transpose(2, 1))                      # (B, N, 64)  
12  
13         feature_transform = TransformationNet(64)(x) # (B, 64, 64)  
14         x = torch.bmm(x, feature_transform)          # (B, N, 64)  
15  
16         x = x.transpose(2, 1)                         # (B, 64, N)  
17         x = F.relu(nn.Conv1d(64, 64, 1)(x))           # (B, 64, N)  
18         x = F.relu(nn.Conv1d(64, 128, 1)(x))          # (B, 128, N)  
19         x = F.relu(nn.Conv1d(128, 1024, 1)(x))         # (B, 1024, N)  
20         x = nn.MaxPool1d(num_points)(x)                # (B, 1024, 1)  
21         x = x.view(-1, 1024)                          # (B, 1024)  
22         return x                                     # (B, 1024)
```

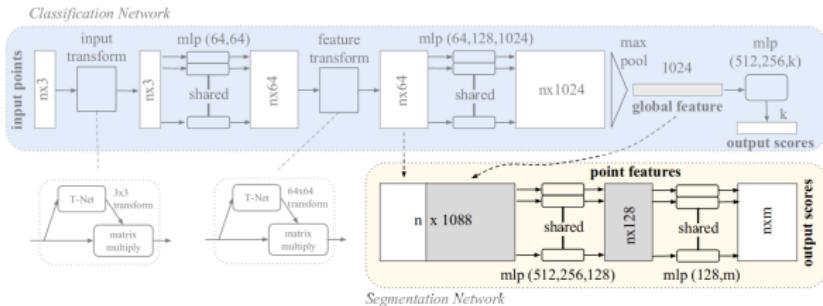
PointNet PyTorch: Classification



Classification Network

```
1 class ClassificationPointNet(nn.Module):  
2  
3     def forward(self, x): # x.shape = (B, N, C)  
4         x = BasePointNet()(x) # (B, 1024)  
5  
6         x = F.relu(nn.Linear(1024, 512)(x)) # (B, 512)  
7         x = F.relu(nn.Linear(512, 256)(x)) # (B, 256)  
8         x = nn.Dropout(0.3)(x) # (B, 256)  
9         x = nn.Linear(256, num_classes)(x) # (B, num_classes)  
10        x = F.log_softmax(x, dim=1) # (B, num_classes)  
11        return x # (B, num_classes)
```

PointNet PyTorch: Segmentation



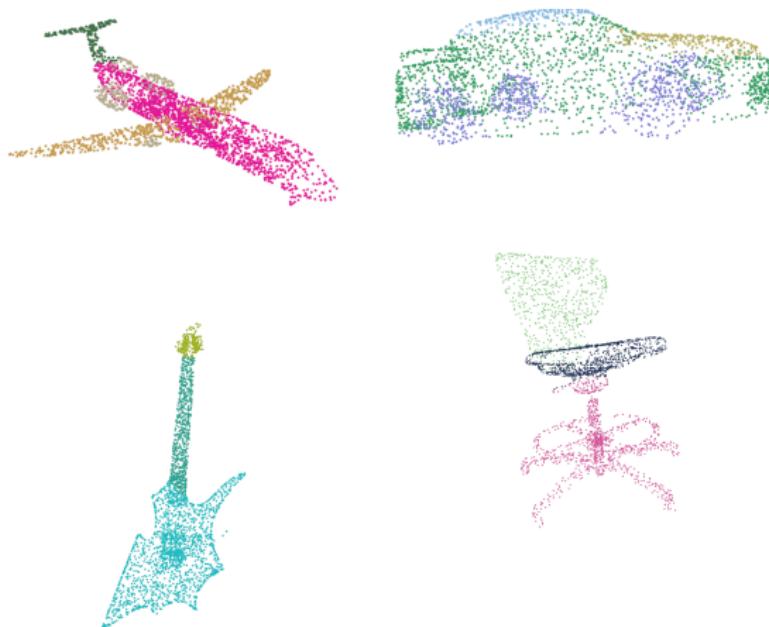
Segmentation Network

```
1 class SegmentationPointNet(nn.Module):  
2  
3     def forward(self, x): # x.shape = (B, N, C)  
4         x = BasePointNet()(x) # (B, N, 1024 + 64)  
5  
6         x = x.transpose(2, 1) # (B, 1024 + 64, N)  
7         x = F.relu(nn.Conv1d(1088, 512, 1)(x))# (B, 512, N)  
8         x = F.relu(nn.Conv1d(512, 256, 1)(x)) # (B, 256, N)  
9         x = F.relu(nn.Conv1d(256, 128, 1)(x)) # (B, 128, N)  
10        x = nn.Conv1d(128, num_classes, 1)(x) # (B, num_classes, N)  
11        x = x.transpose(2, 1) # (B, N, num_classes)  
12        x = F.log_softmax(x, dim=1) # (B, N, n_classes)  
13        return x # (B, N, n_classes)
```

Experiment on ShapeNet

ShapeNet dataset¹: 16 shapes, segmented in parts (up to 6/shape)

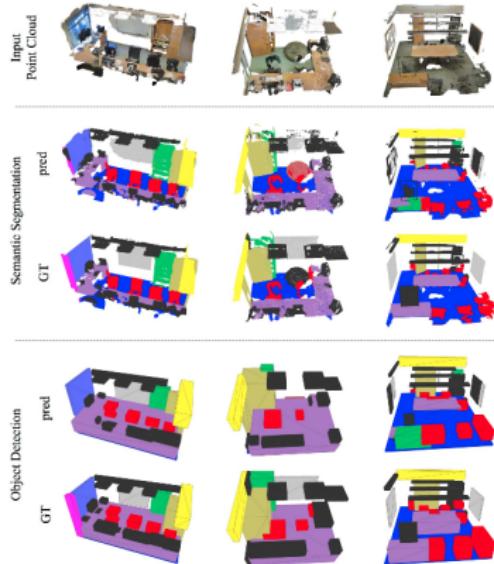
airplane (4027)	earphone (73)	cap (56)	motorbike (336)
wings tail	body headband	peak earphone	wheel light gas tank seat handle
engine	engine	panel	
bag (83)	mug (213)	laptop (452)	table (8420)
handle body	handle	keyboard	top leg
	handle	keypad	top
guitar (793)	knife (420)	rocket (85)	lamp (2308)
body head neck	handle blade	fin body	shade base tube
	blade	nose	
chair (6742)	pistol (307)	car (7496)	skateboard (152)
seat back	handle barrel trigger	wheels hood	deck wheel
arm leg			



¹Source: A Scalable Active Framework for Region Annotation in 3D Shape Collections

Others

Semantic scene segmentation



Credit¹

MNIST classification

	input	error (%)
Multi-layer perceptron [22]	vector	1.60
LeNet5 [12]	image	0.80
Ours PointNet	point set	0.78

Credit¹

¹ "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", Qi et al.

Overview

1 Point Clouds

- Definition
- Visualization
- Learning tasks

2 Learning features from point clouds

- Learning features from regular data
- Point cloud problem statement

3 PointNet

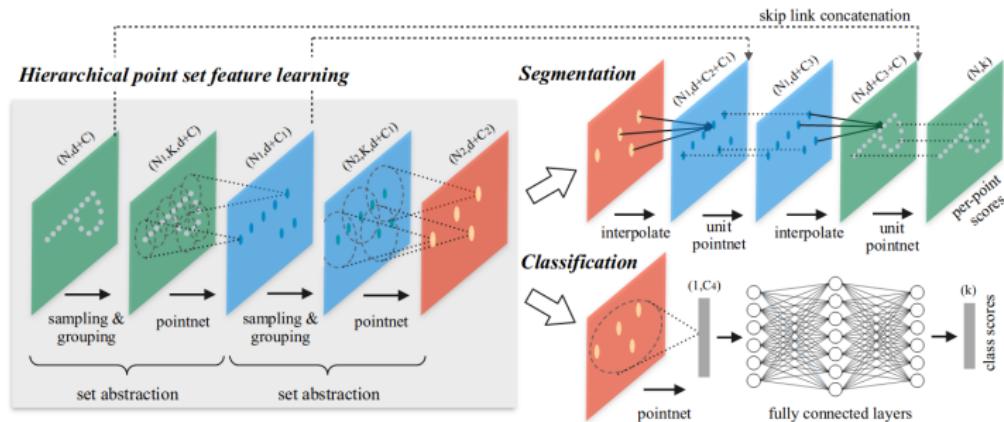
- PointNet proposal
- PointNet architecture
- PointNet PyTorch implementation
- Experiments

4 Beyond PointNet

- PointNet++
- PointCNN

PointNet++

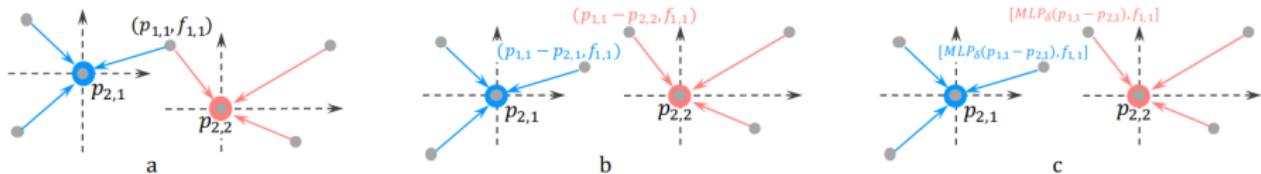
"PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space", Qi et al



Main idea: applying recursively PointNet on a nested partitioning of the input point set (~ stacked CNN layers for 2d images) to capture local structure induced by the metric space.

PointCNN

"PointCNN: Convolution On χ -Transformed Points", Li et al



Main idea: Learn a χ -transformation (local) that transforms features associated with points into a canonical space (\sim invariant to permutation) and then apply a standard convolution.

Thank you for your attention

PyTorch vs TensorFlow

- Dynamic graph (PyTorch) vs static graph (Tensorflow)
- PyTorch is more "pythonic"
- PyTorch is cleaner: better docs, no redundant API...
- Tensorflow has a better ecosystem: e.g. Tensorboard
- Tensorflow is more "production ready"

3D representation

Credit: "Deep Learning Advances on Different 3D Data Representations: A Survey", Ahmed et al

